

Predictive and interpretable models for cancer-specific identification of drug-sensitive tumours from their multi-omics profiles

--Manuscript Draft--

Manuscript Number:		
Full Title:	Predictive and interpretable models for cancer-specific identification of drug-sensitive tumours from their multi-omics profiles	
Article Type:	Research	
Funding Information:	Canceropôle PACA (NA)	Dr. Pedro Ballester
Abstract:	<p>In silico models to predict which tumours will respond to a given drug are necessary for Precision Oncology. However, predictive models are only available for a handful of cases (i.e. drugs acting on tumours of a given cancer type). A way to generate predictive models for additional cases is with suitable machine learning algorithms that are yet to be applied to existing data sets. Here we apply XGBoost allied with a stringent feature selection approach, which is advantageous for these high-dimensional problems. Thus, we identified and validated 118 predictive models for 62 drugs across five cancer types by exploiting four molecular profiles (sequence mutations, copy-number alterations, gene expression and DNA methylation). Predictive models were found in each cancer type and with every molecular profile. On average, no omics profile or cancer type obtained models with higher predictive accuracy than the rest. However, within a given cancer type, some molecular profiles were overrepresented among predictive models. For instance, CNA profiles were predictive in BRCA cell lines, but not in SCLC cell lines where GEX and DNA methylation profiles were most predictive. Lastly, we identified the best XGBoost model per cancer type and analysed their selected features. For each model, some of the genes in the selected list had already been found to be individually linked to the response to that drug, providing additional evidence of the usefulness of these models and the merits of the feature selection scheme.</p>	
Corresponding Author:	Pedro Ballester INSERM FRANCE	
Corresponding Author Secondary Information:		
Corresponding Author's Institution:	INSERM	
Corresponding Author's Secondary Institution:		
First Author:	Stefan Naulaerts	
First Author Secondary Information:		
Order of Authors:	Stefan Naulaerts	
	Michael Menden	
	Pedro Ballester	
Order of Authors Secondary Information:		
Additional Information:		
Question	Response	
Are you submitting this manuscript to a special series or article collection?	No	
Experimental design and statistics	Yes	

<p>Full details of the experimental design and statistical methods used should be given in the Methods section, as detailed in our Minimum Standards Reporting Checklist. Information essential to interpreting the data presented should be made available in the figure legends.</p> <p>Have you included all the information requested in your manuscript?</p>	
<p>Resources</p> <p>A description of all resources used, including antibodies, cell lines, animals and software tools, with enough information to allow them to be uniquely identified, should be included in the Methods section. Authors are strongly encouraged to cite Research Resource Identifiers (RRIDs) for antibodies, model organisms and tools, where possible.</p> <p>Have you included the information requested as detailed in our Minimum Standards Reporting Checklist?</p>	<p>Yes</p>
<p>Availability of data and materials</p> <p>All datasets and code on which the conclusions of the paper rely must be either included in your submission or deposited in publicly available repositories (where available and ethically appropriate), referencing such data using a unique identifier in the references and in the “Availability of Data and Materials” section of your manuscript.</p> <p>Have you have met the above requirement as detailed in our Minimum Standards Reporting Checklist?</p>	<p>Yes</p>

Predictive and interpretable models for cancer-specific identification of drug-sensitive tumours from their multi-omics profiles

Stefan Naulaerts¹, Michael Menden² and Pedro J. Ballester^{1*}

¹ Cancer Research Center of Marseille, INSERM U1068, F-13009 Marseille, France; Institut Paoli-Calmettes, F-13009 Marseille, France; Aix-Marseille Université, F-13284 Marseille, France; and CNRS UMR7258, F-13009 Marseille, France.

² Helmholtz Zentrum München, Germany.

* correspondence to: Pedro Ballester, email: pedro.ballester@inserm.fr

ABSTRACT

In silico models to predict which tumours will respond to a given drug are necessary for Precision Oncology. However, predictive models are only available for a handful of cases (i.e. drugs acting on tumours of a given cancer type). A way to generate predictive models for additional cases is with suitable machine learning algorithms that are yet to be applied to existing data sets. Here we apply XGBoost allied with a stringent feature selection approach, which is advantageous for these high-dimensional problems. Thus, we identified and validated 118 predictive models for 62 drugs across five cancer types by exploiting four molecular profiles (sequence mutations, copy-number alterations, gene expression and DNA methylation). Predictive models were found in each cancer type and with every molecular profile. On average, no omics profile or cancer type obtained models with higher predictive accuracy than the rest. However, within a given cancer type, some molecular profiles were overrepresented among predictive models. For instance, CNA profiles were predictive in BRCA cell lines, but not in SCLC cell lines where GEX and DNA methylation profiles were

most predictive. Lastly, we identified the best XGBoost model per cancer type and analysed their selected features. For each model, some of the genes in the selected list had already been found to be individually linked to the response to that drug, providing additional evidence of the usefulness of these models and the merits of the feature selection scheme.

BACKGROUND

Large-scale cancer *in vitro* pharmacogenomics databases have been generated over the last six years. The most well-known are CCLE[2], CTRP[3] and GDSC[4–6], which provide a deep molecular characterisation (single nucleotide variants, copy number alterations, DNA methylations and gene expression levels) of large panels of cell lines prior to be tested with hundreds of drugs. Those data sets are growing ever since, for example, the latest GDSC study⁵ employed 265 drugs characterised in 1,001 cell lines, with 990 cell line drug responses being readily available. CCLE, CTRP and GDSC have expanded their cell line panels by more than 16 fold compared to the pioneering NCI-60 panel[7], enabling to study over 30 cancer types in detail, albeit restricting their screening effort to a smaller set of clinical and pre-clinical compounds. These screening efforts describe the heterogeneous landscape of cancer and explore pharmacological responses to drug treatments, which ultimately enables Precision Oncology.

The resulting pharmacogenomics databases have in turn spurred the generation of a range of computational models for drug response prediction in cancer cell lines[8–10]. These predictive models have been shown to be accurate in a few binomials drug-cancer type [6,11,12]. Building these models constitutes a strongly high-dimensional problem and, as such, it should employ some form of feature selection. It is indeed error-prone to directly relate many thousands of molecular features to a far smaller number of drug-treated cell lines. Consequently, some studies[6,13] have followed a knowledge-based approach to select likely

relevant subsets of features: e.g. restricting to those genes in the cancer gene census set[14], or to restrict to mutations also observed in TCGA tumour profiles[6]. Purely data-driven feature selection has also been employed to mitigate this form of overfitting, e.g. studies using feature selection embedded in Random Forest (RF)[15–18].

The DREAM 7 challenge[12] attempted to establish which families of algorithms performed best to predict drug sensitivity for a panel of 53 breast cancer cell lines. Although no substantially superior algorithm was found, it was observed that the top-performing models were those accounting for nonlinearities. In that study the best performing model was a Bayesian multitask multiple kernel learning approach, while the runner-up was built upon an underlying RF model. These studies are still scarce in that accurate predictors are still needed for many drugs across many cancer types[19] and comparative benchmarking is even rarer.

One way to identify accurate predictors for orphan drug-cancer type pairs, is to apply new machine learning algorithms. XGBoost (XGB for short)[20] has frequently been part of winner solutions in Kaggle (www.kaggle.com) competitions[20]. Like the well-established RF algorithm, XGB is an ensemble tree-based model built on CART trees. However, XGB adds gradient boosting and is known to require substantially less training time than RF on most problem instances[21]. XGB has also been found to be much faster than deep neural networks while achieving practically the same average performance the related problem of QSAR modelling[21]. In addition, XGB has excelled at predicting drug-target binding affinities[22], where it outperformed widely-used regularised least-squares techniques on the studied data sets. As its constituent tree models ignore by construction irrelevant features, XGB has embedded feature selection, which makes it an algorithm suitable to analyse high-dimensional data.

In this study, we apply XGB algorithms to GDSC data to identify novel cancer-specific models predictive of drug response. In addition to standard XGB, we introduce and apply a

variant integrating an adaptive feature selection scheme. This new scheme is called Optimal Model Complexity (OMC) and it complements the embedded feature selection of XGB by searching for a much smaller subset of features, which often results in less overfitting and thus more predictive models. An additional advantage of the OMC strategy is that a small set of features allows better interpretability and less resource-intensive application to other tumour samples. Here we show the complementarity of both XGB algorithms across the drug-cancer type binomials for which new predictors were found. Furthermore, we identify the most predictive molecular profile for each of these binomials. XGB-OMC predictors exploiting gene expression or DNA methylation profiles are particularly innovative, as these profiles have not been previously analysed with nonlinear regression algorithms able to find a concise subset of predictive features.

DATA DESCRIPTION

We defined each case as a drug and cancer type pair, where cell lines can be described by four omics profiles. In terms of genomic profiles, we analysed SNV and CNA data. Both were encoded as binary features, each feature indicating whether a mutation was detected in the gene or if a chromosomal region had copy number alterations, respectively. On the other hand, gene expression intensities (RMA normalized mapped to gene) and methylation β -values of the CpG islands were used as real-valued features. Regarding drug sensitivity data, we found 990 cell lines with 212,774 log IC₅₀ values for 265 drugs in 30 cancer types in the GDSC (49,576 measurements did not have an associated cancer type and were subsequently discarded). Once we restricted the dataset to the cancer types with at least 45 cell lines, 3168 unique tuples of drug, cancer type and tumour profile type remained. Each of these cases had between 45 and 62 fully-profiled data instances, comprising 219 drugs across five cancer types (BRCA, COADREAD, LUAD, SCLC and SKCM). Further description of the input data and their pre-processing is available in the Methods section.

To be able to assess how predictive are the generated models, we split the cell lines of each case into a training set (used also for model selection via cross-validation) and an external test set. There is hence no overlap between any training and test sets. Every test set contained 10 cell lines and was stratified to ensure that it is representative of drug responses for the considered case. The number of training cell lines, n , varies from case to case. For the training set of each case, OMC considered the top k features, with k ranging from 2 to $n/2$ and n being the number of cell lines in the training set of the considered case. This was done in order to generate a sequence of $(n/2)-1$ nested feature subsets, each leading to a different training set using the same cell lines. We set $n/2$ as the highest model complexity to have at least two training data instances per feature (i.e. two training cell lines per feature as a heuristic to identify models that only require a concise list of features). The OMC was chosen as the top k features resulting in the best median performance over the five folds. The full list of features for the case is also used to build a baseline model with all features and running them on the same training set as the OMC strategy.

In order to select the OMC, features need to be ranked first. We used the p -values obtained from univariate statistical tests to correlate between each feature and the corresponding labels (log IC50 values). For a given case and complexity, p -values between each feature and drug response were calculated using only the training folds (i.e. excluding the validation fold) within the 5-fold stratified cross-validation (see Figure 1). Note that the test set is not used at any step of this training process, in particular the determination of the OMC uses the training set exclusively. For discrete features (SNV, CNA), p -values of the Wilcoxon rank-sum test were used. For continuous features (gene expression intensities, methylation β -values), the correlation coefficient of the Spearman test was used. We opted to use univariate filter methods to avoid long computational times, making assumptions or biasing predictive performance, as previously reported in literature[23] (this study found univariate filters to be

among the top-performing feature selection algorithms in a related high-dimensional problem).

We estimated the OMC using a five-fold stratified cross-validation for model selection. To this end, log IC50 values were randomly distributed to one of five folds, yet their ranking within a fold was maintained from small to large in order to represent a fold representative of the test set and the full drug response range. Sizes of validation set folds varied between 7 and 11 cell lines, which is comparable to the size of the external test set. We therefore expected the validation folds to provide a good representation of the behaviour of the model in the test set. To guarantee that the OMC generalises well, we required that the minimal Spearman correlation (R_s) of all five validation folds exceeded a R_s of 0.25 of all five validation folds exceeded a R_s of 0.25 (this is a cut-off previously used to model these data sets[6]). Once the OMC is determined, p-values are calculated using the entire training set and only the top k_{opt} features are employed to train the XGB-OMC model, which is applied to predict drug sensitivities in the test set (see Figure 1).

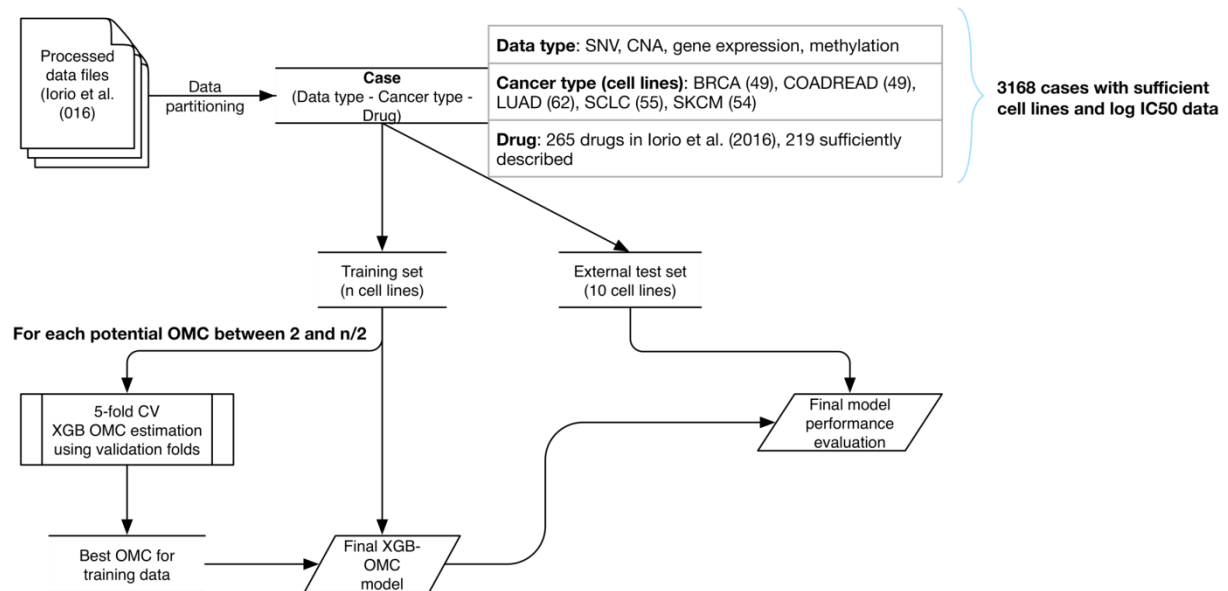


Figure 1: Schematic outline of the procedure to determine the Optimal Model Complexity (OMC) for each case and evaluating the resulting models. Data pre-processed in the latest

GDSC study[6] is partitioned into unique cases, then split into training and test set. A five-fold cross-validation using stratified folds is used to establish which are the top k features (where k takes all the integer values between 2 and $n/2$, with n being the number of training cell lines) leading to the most predictive XGB regression model over the five cross-validation runs. Once the optimal k (k_{opt}) has been determined, the top k_{opt} features on the entire training set are identified. This is repeated for each case, giving rise to the predictive models discovered and analysed in this study.

ANALYSES

Building predictive cancer-specific models

For a model to be regarded as potentially predictive, there needs to be sufficient correlation between its predictions and the true observed log IC50 values. We thus required that a minimum R_s of 0.25 in every validation fold for a model to be considered as potentially predictive. For a potentially predictive model to be confirmed as predictive, we also required the model needed to have R_s greater than 0.25 in the test set. By eliminating those with performance below this test set threshold, the number of models was reduced from 268 to 77 (28.73%). The 77 predictive models cover 62 unique drugs across all four data types in the five cancer types with sufficient cell lines. These models are specified in the supplementary material (Supplementary file 1).

To further assess whether these 77 models are truly predictive, we tested each selected OMC subset by training models for a five-fold cross-validation on the full dataset (i.e. merging training and test sets for the corresponding cases), using 10 random seeds. In Figure 2, we compared the performances of XGB-OMC models with those from models trained using case-specific randomly permuted log IC50 values. As can be seen from the plots, even with just 10 random seeds, there is a substantial difference between the models trained with data (blue) and those trained using randomly permuted log IC50s (green). Comparison using a

1 paired t-test for Spearman correlations from the predictive model and their corresponding
2
3 permutated variant indicated a highly significant difference between both model types (p-
4
5 value $< 1.6 \cdot 10^{-16}$ across cases), indicating that these model predictions are extremely unlikely
6
7 to be due chance. As can be seen in Figure 2, the XGB-OMC models generally outperform
8
9 models trained on randomly permutated data (note much higher Spearman correlations in the
10
11 test set over all 10 seeds). This shows that the OMC models capture properties relevant to the
12
13 drug response. Only one model was found to be slightly problematic when compared to its
14
15 permutated counterpart (the DNA methylation model for COAD/READ cell lines treated
16
17 with PFI-1). Compared to most well-predicted cases, the variance in log IC50 for PFI-1 on
18
19 the relatively low end (0.209 for PFI-1) when compared to the variances for well-predicted
20
21 drug-cancer combinations in the validation and test sets, such as Thapsigargin in SCLC
22
23 (0.588), Bleomycin in BRCA (0.653), Mitomycin C in LUAD (0.473), PI-103 in
24
25 COAD/READ (0.776). This difference was much less pronounced when compared to the best
26
27 SKCM case (Pazopanib, log IC50 variance of 0.268).
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65

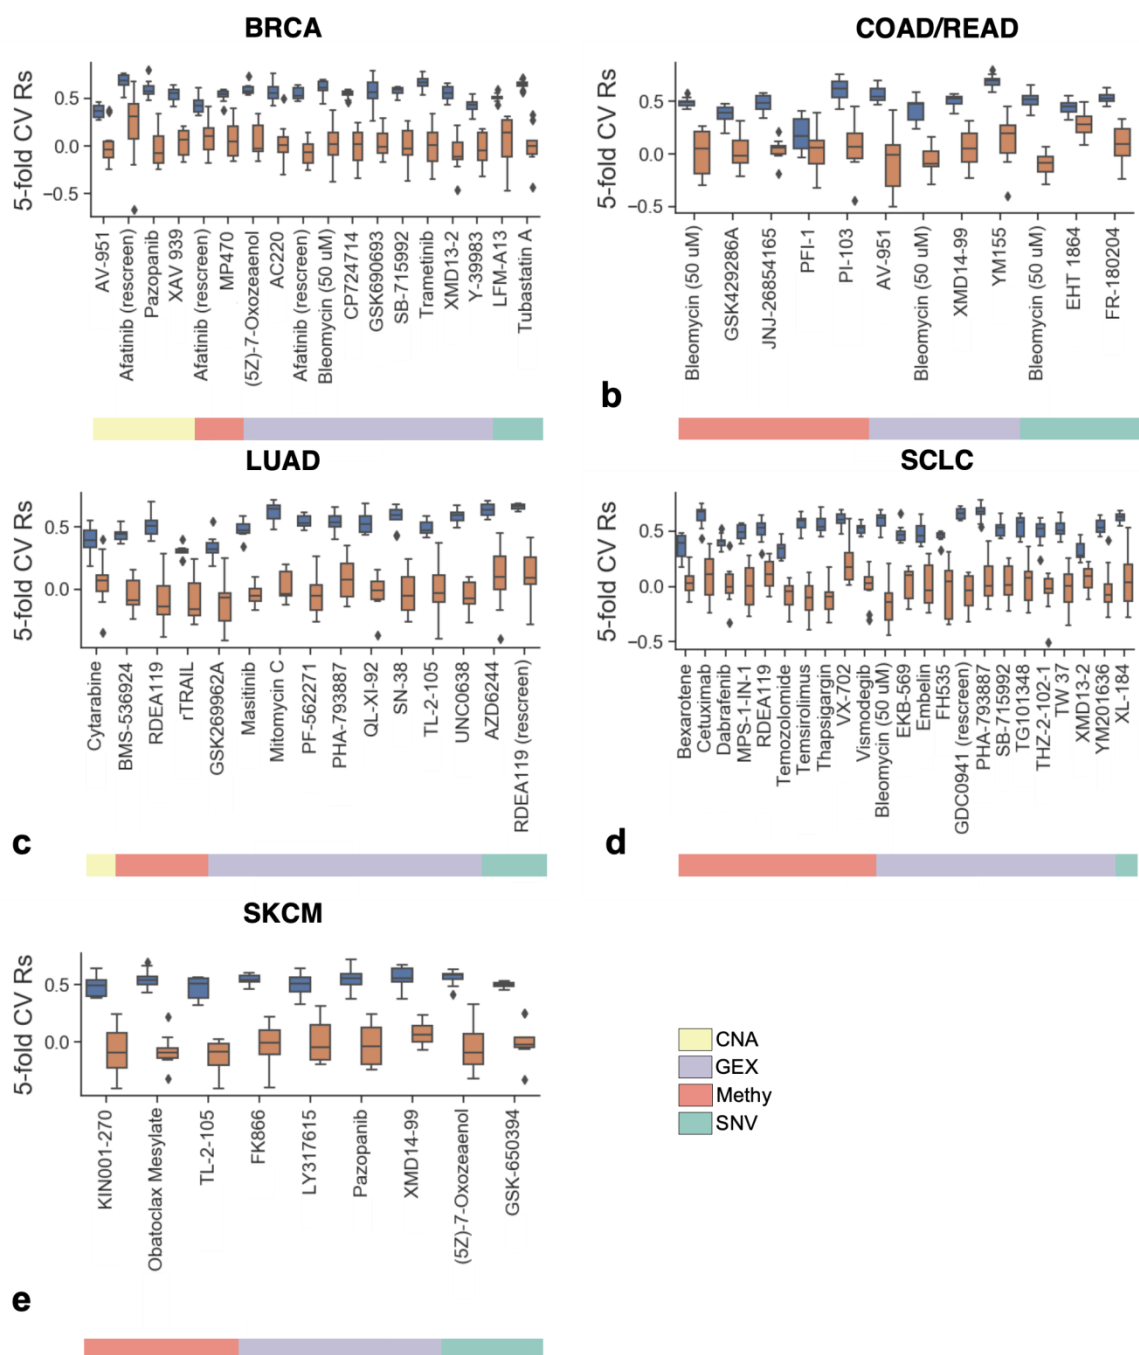


Figure 2: Performances of predictive XGB-OMC models (blue) and their corresponding models using permuted log IC₅₀ values (orange). In each case, performance is calculated as Spearman's correlation between predicted and observed log IC₅₀ values (the boxplot is formed with the five validation R_s from the 5-fold cross-validation using the full dataset for that case). All five cancer types are shown (a: BRCA, b: COADREAD, c: LUAD, d: SCLC, e: SKCM), each predictive model identified by its drug (name) and molecular profile (horizontal colour bar below each plot; see legend for correspondence). Of all models, the

model for PFI-1 in COAD/READ is by far the worst performer when compared to the permuted log IC50s (p-value=0.27). This model obtains borderline performance on not only the test set (R_s of 0.31), also on the validation set (R_s of 0.29). The odds that the other models are being obtained by chance are very unlikely ($p < 2.2 \cdot 10^{-16}$, t-test).

For all 268 cases that were potentially predictive (R_s greater than 0.25 in all validation folds), we compared the models using the OMC with the corresponding models using all features. As previously seen (Figure 2), 77 of the 268 models were predictive in our external test set when using the OMC. Interestingly, we found that 96 models of these 268 were predictive using all features. Therefore, using the XGB-all results in a slightly larger number of predictive models (35.82%) when compared to XGB-OMC (28.73%). Of the 96 XGB-all cases that were predictive, 55 were also predictive using XGB-OMC, thus indicating consistency between both approaches. In 24 of the 55 shared models, XGB-OMC outperformed XGB-all, with the latter seemingly more predictive. However, we found that the average difference in test set performance between the two was 0.022 in favour of XGB-all, a not statistically significant difference (paired t-test; $p=0.374$). We cannot therefore conclude that both methods perform dissimilarly in these 55 cases in common. It is important to note that XGB-OMC resulted in a major reduction of the number of features: 84% of the considered features were on average removed without losing much, if any, predictive performance compared to XGB-all.

Figure 3 shows a high degree of complementarity between XGB methods. For 41 of the 96 predictive XGB-all models, the OMC pre-ranking fails to capture potentially predictive characteristics. By contrast, XGB-OMC is predictive in 22 cases where XGB with all features was not. Therefore, the use of both methods is required to obtain predictive models for a larger number of drugs. For those cases below but close to the diagonal, the use the XGB-OMC models is preferred to reduce the feature space with negligible loss of performance.

1 From a practical perspective, it is interesting to see that for over half the cases for which a
2 prediction using the standard all features predictions is possible, these can just as well be
3 described with the OMC subset, as hypothesis testing can then be performed in targeted
4 research experiments, while avoiding the need for costly exome- or genome-wide
5 experiments[27] and qualified personnel that may not be required.
6
7
8
9
10

11 Also of interest are the cases for which XGB with all features did generate less predictive
12 models. In each plot in Figure 3, these cases are shown above the dashed diagonal line. It is
13 therefore important to consider using the OMC approach as a complement to traditional
14 machine learning, as doing so will add 22 cases that were poorly predicted using the same
15 dataset. In Figure 3, all models that could be considered predictive using either the OMC or
16 the all features method are shown. In total, our study has identified 118 novel XGB models
17 with predictive value (77 predictive with XGB-OMC, complemented by 41 XGB-all), all of
18 which are provided together with their OMC feature set in Supplementary file 1. These 118
19 feature lists and their corresponding models constitute a large set of complex hypothesis of
20 how cancer cell lines respond to a drug.
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65

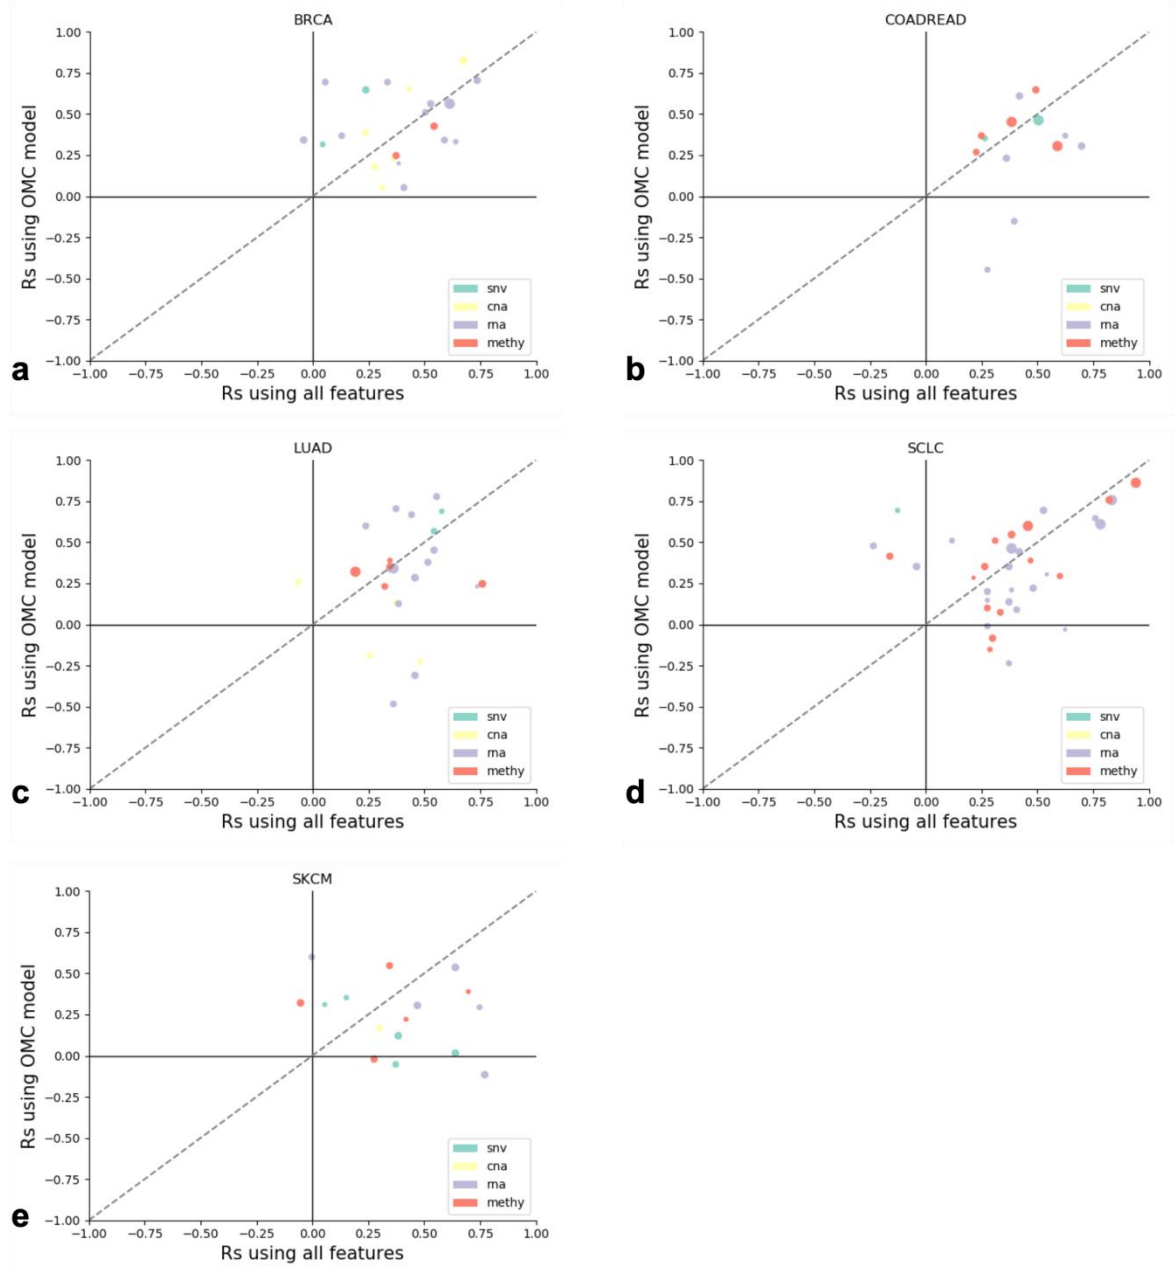


Figure 3: Test set Spearman's correlations across all 118 cases with predictive models (i.e. R_s in validation folds > 0.25 and R_s in test set > 0.25). The case with the most predictive omics profile is shown as the node colour, while the size of the dots represents the OMC. The dashed grey line represents a 1:1 ratio between both. In each of the five cancer types (a: BRCA, b: COADREAD, c: LUAD, d: SCLC, e: SKCM), results from both methods are remarkably different overall, despite using the same base learner (XGB) as well as the same set of training and test sets. This strong difference is evidence of their complementarity.

Identifying the most predictive molecular profile of each drug and cancer type

We observed that the identity of the most predictive profile depends strongly on the considered case. Note that there are more predictive models than drug-cancer type binomials, as some of these cases are well predicted using more than one profile. For example, CNA profiles were employed by only 13 of the 118 predictive models (11 of these 13 were the best model for the case). Surprisingly, we found that only 13 SNV models were predictive (12 of these were the best model for the case). This was in stark contrast to real-valued profiles, that did indeed seem to be much more informative: 60 predictive models using gene expression data (56 best models for the case) and 32 predictive models DNA methylation data (31 best model for the case). Although much information for the binary data types (SNV, CNA) may be lost in the feature creation step (e.g. choice of only census genes, disregarding the specific mutation, disregarding the number of copies, ...), our results based on the GDSC features seem to suggest that for exploratory analysis, gene expression may be the most likely data type to result in predictive models (50.8% of predictive models). Test set performances by means of Spearman correlation are shown in Figure 4 for both molecular profile types (left) and cancer types (right).

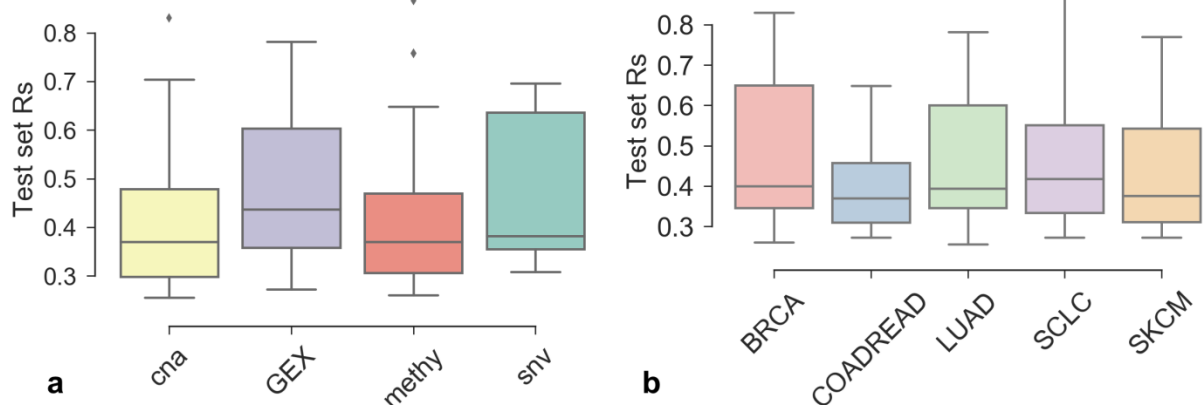


Figure 4: Test set performances of the 118 models identified as predictive in this study. We grouped all test set performances of the predictive models per omics profile (Panel A). Interestingly, there is no universally better omics data type. However, gene expression (GEX)

was characterized by high predictive performance and it resulted in the largest number of predictive models. In panel B (right), we visualize the predictive performance of all 118 XGB-OMC and XGB-all models when grouped per cancer type. As in panel A, no cancer type is significantly better predicted than another. Both plots show that predictive models can be built for each omics profile and cancer type.

In Figure 4B, we can see that the test set performances for models belonging to the different cancer types. These differences are not statistically significant (p-values > 0.05, t-test). In total, we established 24 predictive models for BRCA, 16 for COADREAD, 25 for LUAD, 37 for SCLC and 16 for SKCM across all data types. In contrast to omics profiles, the distribution of the 118 predictive models across cancer types is much more even, complementing their equivalent overall performance characteristics, which suggest we are not biasing towards a single cancer type. The two cancer types with the largest number of cell lines (LUAD and SCLC) are also those with the largest number of predictive models, which further supports the importance of larger data sets. An ANOVA analysis performed in R could not find a statistically significant effect of the cancer type, indicating that no cancer type was better predicted than the rest.

In the majority of cancer types, gene expression results in substantially more predictive models than the other three molecular profiles. This is not surprising, given that overall 60 of the 118 models exploited these data sets. However, the relative numbers of predictive models across omics profiles are strongly cancer-specific, as it can be seen in Figure 5. Interestingly, some cases, such as Afatinib and Bleomycin on BRCA, have predictive models with more than one molecular profile (Figure 5 only presents test set performances for the best profile, which are those with the highest test set Rs). This figure also shows the drugs with predictive models on each cancer type (BRCA: 22 drugs, COADREAD: 13 drugs, SKCM: 16 drugs, LUAD: 24 drugs, SCLC: 35 drugs).

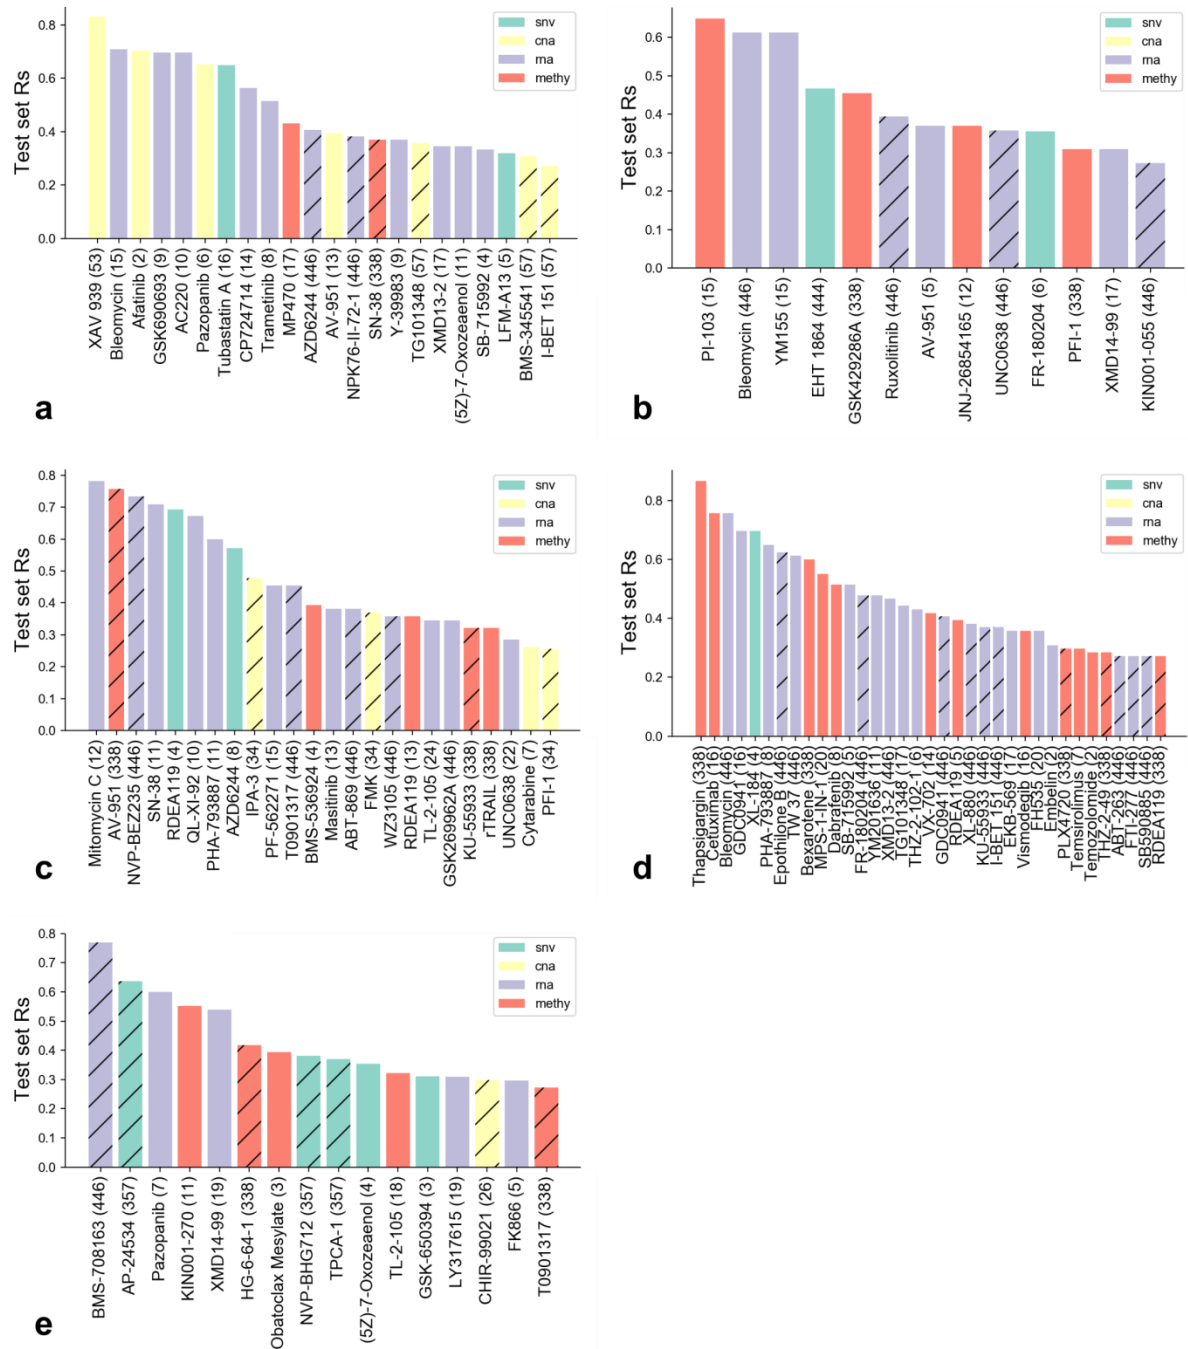


Figure 5: Test set Spearman's correlation (Rs) between predicted and observed logIC50 of the best cases from each cancer type (a: BRCA, b: COADREAD, c: LUAD, d: SCLC, e: SKCM). The data type of the best case (i.e. highest Rs) is shown as colour. If the bar is hatched, this indicates that using the XGB-all led to a predictive model where XGB-OMC did not (test set Rs<0.25). Note that gene expression is the single-most informative profile across

the considered cancer types. The contributions the other omics data types can vary greatly based on the cancer type.

It is noteworthy that of the 12 cases best predicted with CNA features, 7 are found in BRCA. Similarly, SCLC seems to have a proportionally larger presence of drugs that can be predicted with methylation information (13 of the 35 models), while CNA does not seem to play as much of a role as it does in BRCA. On the other hand, SNV seems to be more relevant for SKCM. Although this relation between predictive omics data types and cancer type depends strongly on the drugs tested, extensive reviews have been made that describe the importance of a given data type to specific cancers, such as the role of DNA methylation in SCLC[28].

The best XGB-OMC model per cancer type as case studies

As case studies for discussion, we selected the best XGB-OMC model per cancer type, resulting in five distinct case studies consisting of two gene expression models, two DNA-methylation models and a CNA model. These models, together with test set performance of both corresponding XGB-OMC and XGB-all are shown in table 1. The R^2 metric (coefficient of determination), which is an alternative parameter to gauge model performance, is also shown. Two models use all features, while the other three require far fewer features to be predictive. Note that the Pazopanib-SKCM could not be predicted at all using all 450 GEX features, unlike when using the OMC to build XGB models with just seven of them.

Profile	Cancer type	Drug Name	OMC	Test R_s (XGB-OMC)	Test R_s (XGB-all)	Test R^2 (XGB-OMC)	Test R^2 (XGB-all)
Methy	SCLC	Thapsigargin	338 (all)	0.867	0.939	0.535	0.609
CNA	BRCA	XAV 939	53 (all)	0.830	0.673	0.347	0.382
GEX	LUAD	Mitomycin C	12	0.782	0.552	0.465	0.280
Methy	COADREAD	PI-103	15	0.648	0.491	0.395	0.296
GEX	SKCM	Pazopanib	7	0.600	-0.006	0.321	-0.231

Table 1: The most predictive XGB-OMC models per cancer type. Optimal complexity, test set performance (Spearman) and the test set performance (Spearman) of the corresponding all features models are shown (BRCA: breast invasive carcinoma, COADREAD: Colon adenocarcinoma/Rectum adenocarcinoma, LUAD: Lung adenocarcinoma, SCLC: Small cell lung cancer, SKCM: Skin cutaneous melanoma). For each of these models, validation set performance exceeded $R_s > 0.25$ in all validation folds.

We highlighted the best XGB-OMC model per cancer type to discuss in detail (Table 1). The first of these five case studies is the DNA methylation model to predict the response of SCLC cell lines to Thapsigargin. The plant-derived compound thapsigargin is a Ca^{2+} -ATPase antagonist that can increase intracellular Ca^{2+} concentrations in small cell lung cancer cells and thus has an effect on excreted autocrine growth factors, such as serotonin[29], which is represented at the gene expression level[30]. Furthermore, SCLC shows very distinct methylation patterns that strongly correlate with E2F expression. Activation of EZH2 has been shown to be a direct consequence of RB1/E2F pathway dysregulation in SCLC[31]. EZH2 is a histone methyltransferase, with EZH2 inhibiting tumour growth in PDX models[28]. Our results, based on cell lines, are consistent with findings found in the literature, which indicate that the RB1/E2F pathway, and consequently this DNA methylation pattern, may hold predictive value for SCLC treatment strategies.

The second case study is the copy-number alteration model to predict the response of BRCA cell lines to XAV 939. XAV-939 inhibits tankyrase 1 and 2, thus inhibiting Wnt/ β -catenin-mediated transcription and stabilizes the cytoplasmic axin levels. Genes associated with the Wnt pathway have high expression in triple-negative breast cancer[32], which makes it an interesting therapeutic prospect to fight this cancer type, which represents between 10-20% of all breast cancer cases. In our feature set, we found all 53 BRCA-specific CNA to be required for the optimal XGB model (24 were amplifications and 29 constituted deletions). We

mapped the CNA-altered region to the gene identifiers of the contained genes with the GDSC-provided files and performed an overrepresentation analysis with pathway terms. Enrichment analysis using KEGG annotation indicated a slight enrichment of BRCA (p=0.029, ID: hsa05224), thanks to presence of APC2, a WNT signaling pathway regulator, E2F3 (also involved in Wnt/ β -catenin regulation), WNT5B (Wnt family member 5), EGFR (epidermal growth factor) and a large number of fibroblast growth factors being present. The presence of several Wnt signaling-related features in the dataset and the model's performance in a cancer type in which the drug has been found to have an effect, increases confidence in the complex hypothesis embodied by the XGB-OMC model.

Gene expression model to predict the response of LUAD cell lines to Mitomycin C. This case considered Mitomycin in LUAD, which again seems to be in a clinically relevant tissue for the drug. Thus far, 11 clinical trials in clinicaltrials.gov have test either completed or ongoing that include the use of mitomycin for lung cancer, it is classified as a structured indication for lung cancers in Drugbank (Drugbank ID: DB00305) and has been used since 1984 for non-small cell lung cancers[33]. Mitomycin is an antibiotic produced by amongst others the actinobacteria *Streptomyces caespitosus*. Our OMC subset contained 12 gene expression features: NCOR2_expr, MCM3_expr, AMOT_expr, POLR2B_expr, ARFGEF2_expr, U2AF1_expr, NDRG1_expr, MYD88_expr, NTRK2_expr, ATF1_expr, SHMT1_expr and HNRPD_L_expr. Interestingly, MYD88 works in close relation with TNF[34], which has been frequently linked to Mitomycin and its Fas/FasL dependent apoptosis in cervical carcinoma[35]. Furthermore, several of the features in the OMC set are associated at the protein level with NQO1, which has been described as having an effect on the sensitivity to Mitomycin C[36]. Finally, protein levels of HNRP proteins, one of which is in the OMC set, are known to rise after DNA damage[37], such as the damage caused by oxidative stress induced by Mitomycin C.

The fourth case study is DNA methylation model to predict the response of COADREAD cell lines to PI-103. PI-103 is an inhibitor of PI3K (Phospho-Inoside 3-kinase) and mTOR. Previous research indicated that the effects of PI-103, such as downregulation of choline kinase α , were observable in prostate and colon carcinoma cell lines using magnetic resonance spectroscopy[38]. PI-103 is known to enhance radiosensitivity in several cell lines (including colorectal) and reduced phosphorylation of AKT1 at serine 473[39]. In addition, PI-103 has also been found to slow down tumour growth in PDX models[40]. Hence, our models identify again a predictive case, for which this drug and cancer type are studied together as potentially interesting in literature. Since PI3K is has a downregulatory effect of the PI3K pathway, we expected some members of our OMC subset to be related to this pathway. We found that a complexity of 15 iCPGs best describes the drug response to PI-103. Interestingly, several of these were also present in other cancer types in which the effect of PI-103 has been described. We used the GDSC-provided file (TableS2H.xlsx) to obtain the names of the genes regulated by the considered iCPGs. This resulted in the following set of 18 genes: RAPGEFL1, RRS1, ADHFE1, LIPG, OXR1, ADPRH, ENC1, GSTT1, HSD11B2, ZNF3, DNAH10, PAQR8, SIAH1, HOXA11AS, HOXA11, KAT2A, HSPB9 and P4HB. P4HB is closely associated with AKT1, which is known to be affected by PI-103[39]. Moreover, GSTT1, HOXA11 and SIAH1 are associated with TP53 in the StringDB[41] network, the latter which was shown to be affected by PI-103 in AML[42] and is closely associated with AKT1. These studies further supports that the methylation status of the promoter regions of at least some OMC-selected genes are highly relevant for PI-103's drug response.

The last case study is the gene expression model to predict the response of SKCM cell lines to Pazopanib. Pazopanib is a tyrosine kinase inhibitor, with antineoplastic activity. It has been used to treat renal cell cancer and soft tissue sarcoma. The drug acts against vascular

endothelial growth factors 1, 2 and 3, as well as platelet-derived growth factor receptor beta[43]. We found that an OMC of 7 was sufficient for a gene-expression model to be predictive. Our features were TRERF1_expr, KDR_expr, VHL_expr, PIP5K1A_expr, FMR1_expr, PTPRU_expr and EPHA4_expr. None of these are a known target of Pazopanib. However, we again found close associations between members of our OMC subset with the known drug targets. For example, PTPRU is a direct interactor with KIT, with a combined String score of 0.947, indicating very high confidence. KDR associates with the known targets PDGFB (0.559), VEGFR1(0.963) and VEGFR3 (0.943) and is a connector node in the between PDGFB and VEGFR1 and VEGFR2. VHL, also in our OMC set, associates with both KDR and FLT1. Other than finding immediate connections in a protein-protein interaction network, we could find that a clinical trial involving Pazopanib for stage IV cutaneous melanoma was completed in 2010 and reported one patient of the 13 tested with a partial or complete response (<https://clinicaltrials.gov/ct2/show/results/NCT00861913>). It was also further tested in combination with paclitaxel for advanced melanoma, with the authors describing the potential as interesting[44]. Lastly, Pazopanib also had a reported activity in melanoma xenografts[45]. Overall, these observations again suggest that the predicted indication may be clinically relevant.

DISCUSSION

Cell lines have demonstrated to be highly useful cancer models, especially in a pharmacogenomics context. They are cost-efficient, quick to grow and amenable to high-throughput experiments[52]. However, like all disease models, they also suffer from several inherent limitations. For example, intra-tumour heterogeneity, extracellular environment and immune system response are not modelled in cell lines, which may lead to alternative disease models, such as PDX[53], being more predictive. Cell lines are furthermore prone to

divergence across passages[54]. Despite these issues, several cell lines have proven to be relevant by retaining sufficient correlation to primary tumours[52,54,55]. And given their suitability for high-throughput experiments[52], they are also the model for which most phenotypic and pharmacological information is publicly available. For many drugs, only much smaller amounts of pharmacogenomics data sets, if any, has been released using more patient-relevant *in vivo* models[56]. Consequently, often the only data that can be used to predict drug sensitivity comes from cell lines, thus having their niche in guiding precision oncology efforts[52].

Our results indicate that XGB, combined with a strategy searching for the Optimal Model Complexity (OMC), can generate very predictive models using only concise sets of tumour features. The latter means that only a fraction of the initially considered features need to be determined in additional cell lines, which is far less resource-intensive. We further demonstrate that the predictive models are very unlikely to arise by chance (Figure 2). In addition, they are largely complementary to those generated by XGB using all features (Figure 3). Predictive models were found in each cancer type and with every molecular profile (Figure 4). On average, no omics profile or cancer type obtained models with higher predictive accuracy than the rest. However, within a given cancer type, some molecular profiles were overrepresented among predictive models (Figure 5). For instance, CNA profiles were predictive in BRCA cell lines, but not in SCLC cell lines where GEX and DNA methylation profiles were most predictive.

Although we demonstrate the advantages of the approach using only the best-performing models for each cancer type, many more interesting cases were also found (see Figure 5). For example, cases recapitulating the role of KRAS and BRAF, who are extensively used as clinical biomarkers⁴⁸. In LUAD, mutation of KRAS appears by estimation in 17% of all cases, while only 2% carries a BRAF mutation⁴⁸. Given all evidence outlined in their review,

the authors suggest screening BRAF, KRAS, NF1 and gain information on the effective use on anti-EGFR treatment. KRAS was suggested as biomarker for treatment with AZD6244 (Selumetinib), a MEK1⁴⁹, while a more phase II clinical study identified that NF1 was also indicative of tumour response to AZD6244⁵⁰. Our results support the use of mutations in both KRAS and NF1 in combination with those in the other selected genes (KRAS_mut MYH14_mut FN1_mut MAGI2_mut CHD9_mut ARHGEF6_mut NF1_mut EFTUD2_mut) as a potentially viable biomarker set for AZD6244 response. This set is much richer than predictions involving only KRAS and NF1.

POTENTIAL IMPLICATIONS

Given that all of our five cases were found to be in relevant cancer types and showed tight integration with verified or suggested drug mode of action, we believe that the remaining 113 predictive models, as well as the OMC methodology, could also provide useful starting points for researchers investigating how the response of these drugs are affected by the molecular background of tumours.

METHODS

GDSC

All data sets used were made publicly available by the authors of the 2016 GDSC study through their publication-specific download portal, which can be found at http://www.cancerrxgene.org/gdsc1000/GDSC1000_WebResources/Home.html. This page provides several Excel sheets referred to in the Iorio et al paper[6] as supplementary material. All data sets containing cell line phenotypic data, drug response and information regarding primary tumours used in this study, were downloaded through this web portal. The pre-processed data sets were used, as the steps taken by Iorio et al[6] result in a well-annotated list of features with known functional role in cancer.

Cell line information (metadata)

Throughout the GDSC, both the actual name of the cell line, as well as its COSMIC identifier were used. We created a simple mapping between these two identifiers using the Excel sheet named TableS1E.xlsx (tab “TableS1E-CellLines”). This file also contained further information, such as GDSC tissue type, growth medium, growth properties, microsatellite instability and TCGA abbreviation. For each cell line, we only retained the identifiers (COSMIC identifier, sample name) and the TCGA abbreviation of the corresponding cancer type.

Drug response (continuous variable)

Drug responses were provided as the natural logarithm of half the maximal inhibitory concentration in a data matrix (TableS4A.xlsx, tab “TableS4A-IC50s”), which we subsequently converted to logarithms with base 10. Cell line information was then merged with the drug response and other phenotypic properties (genomics, transcriptomics, methylomics). Pre-processing steps for the four molecular profile types (SNV, CNA, gene expression, methylation) are further described below.

Single-nucleotide variants (SNV, binary feature)

Using whole-exome sequencing data for 48 studies with matched tumour-normal samples, Iorio et al[6] combined the outputs of three variant calling algorithms to predict 470 important cancer genes in primary tumours (TableS2A.xlsx, tab “TableS2A-CancerGenes”). In TableS2B.xlsx (tab “TableS2B-TumourVariants”), also available from the download portal, they cover the mutations of all 18,406 genes in primary tumours, while “TableS2C.xlsx” (tab “TableS2C-CellLineVariants”) describes 486,243 variants and their classification (missense, nonsense, frameshift, ess_splice, inframe, stop_lost, gene_fusion) found in 19,100 genes for the 1001 GDSC cell lines in 30 cancer types. 470 gene names appeared in the primary tumour list (“TableS2B.xlsx”), while these 470 predicted cancer driver genes (“TableS2A.xlsx”) were also present in the cell line variant list

(“TableS2C.xlsx”). The presence of mutations in these 470 cancer genes in both cell lines and primary tumours signifies their potential clinical relevance. We assigned a 1 to all cell lines in which the feature (gene) carried a mutation, while zeroes indicated wild type status. This resulted in the 470 SNV features employed.

Iorio et al[6] further performed a driver mutation analysis with 358 driver mutations reported in primary tumours and 310 in cell lines (293 overlapping between both, available as BEM (binary event matrix) files from the download portal). This constitutes a substantial reduction in feature space. We did not make this distinction between driver and passenger mutations, as the relation between the state of a gene mutation (driver, passenger) and the actual drug response is currently not well understood. We therefore opted to err on the side of caution and relied on the dataset itself to identify lists of features with an effect on predictive performance.

Copy number alterations (CNA, binary feature)

CNA data for cell lines was obtained from TableS2G.xlsx (tab “TableS2G-CellLinesRACSS_CNA”). CNAs were identified by Iorio et al[6] through ADMIRE[57] analysis over 27 cancer types. Note that no CNA data was released for SCLC. From TableS2G.xlsx, we extracted all segments listed under the column header “Region identifier”. These segments were annotated with the prefix cna, the cancer type and a number (e.g. cnaLUAD1). Contained genes are also indicated between brackets where applicable (e.g. cnaLUAD35 (MECOM)). In total, we obtained 558 unique segments for cell lines. From TableS2F.xlsx, which contains the altered segments in primary tumours, we again filtered out everything but the contents of the “Region identifier” column. This resulted in 851 altered segments for primary tumours. We again selected the intersection between the data for the cell lines and the primary tumours as the section being most relevant for further studies in patients. All 558 segments in the cell lines appeared in primary tumours, while the other 293

were unique to the primary tumours, as reported in the original article[6]. Similar to the genes carrying SNVs, we opted to indicate a copy number altered state as a binary value (either 1 if True or 0 if not). We used all 558 segments that were not unique to cell lines. These are cancer-specific and hence each cancer type employed a different set of CNA features.

Transcriptome (GEX, continuous feature).

RMA-normalized pre-processed expression information (sanger1018_brainarray_ensemblgene_rma.txt) was also available throughout the article download portal. However, in this case it links out to the main GDSC portal. We used the latest dataset version of 2 March 2017. The file contains the Ensembl gene identifiers for the genes and COSMIC identifiers for the cell lines. Out of the 1018 cell lines in the file, we retained only those with a COSMIC identifier part of the 1001 cell lines presented in the original GDSC study[6]. The list of genes we reduced from the 17,737 Ensembl gene identifiers to the list of 466 genes to represent the information also covered in SNV. To this end, we mapped the identifiers to the HGNC gene name using the mapping files provided by Ensembl BioMart[58] (Supplementary file 2). 446 out of 470 HGNC gene symbols could be successfully mapped to their corresponding Ensembl identifiers. The 24 missing entries could all be attributed to use of a secondary name or lack of presence on the array, which we corrected accordingly. Subsequently, 450 features had both expression and mutation information available and these were used as gene expression features. We retained all gene expression values as continuous variable in order to keep all information content of the data type. Therefore, each cell line was characterized by 450 GEX features.

Methylome (Methy, continuous feature).

The list of informative CpG islands (iCpGs) hypermethylated in primary tumours (“Table S2I.xlsx”, GDSC article-specific download portal) contained 357 unique sites in the column “HyperMethylated iCpG”. We further downloaded the iCpGs found in cell lines by Iorio et

al[6] (TableS2J.xlsx), which contained 338 unique hypermethylated iCpGs and created the intersection between both lists for clinical relevance. All 338 iCpGs found in cell lines occurred in the primary tumours. As the average pre-processed β -values for each of these sites were also available ("F2_METH_CELL_Data.txt"), we opted to keep the content type as rich as possible and use methylation as a continuous feature based on the average β -values per site. We created the methylation data matrix with 338 iCpGs as features and cell lines as index using the mapping file "methSampleId_2_cosmicIds.xlsx", also available from the GDSC study[6].

Generation of non-overlapping training and test sets.

Starting from the pre-processed data, we compiled a data set for each cancer type with at least 45 treated cell lines annotated with response to a given drug. This modelling choice takes as reference the lowest training set size successfully used in the literature[12], where 28 drugs with 35 profiled cell lines as training set. To be on the safe side, we only considered cases with at least 35 training cell lines. Since we also required an external test set of at least 10 cell lines, only cancer types with at least 45 cell lines in the GDSC[4] were analysed.

Per case (i.e. a combination of drug, cancer type and omics data type), an external test containing 10 cell lines was created. Since we intended to create a test set representative of the full drug response range, we ranked all log IC50 values in ascending order and place the second most extreme values (i.e. the log IC50 just below maximal and that just above minimal log IC50 for the considered case) to the external test set. The eight other test set cell lines were selected in an equidistant manner, based on the rank of their log IC50 value. This stratification procedure covers a representative range of the drug response and avoids the situation that may appear with fully random selection in which either all high or all low log IC50s are selected, resulting in a biased test set. These ten cell lines were excluded from the

model training set and were therefore not used at all for model building or selection. The extreme IC50 values are thus always part of the training set.

Optimal Model Complexity estimation.

XGB will also be employed in conjunction with a strategy to identify the Optimal Model Complexity (OMC) for each case (i.e. each trinomial drug-cancer type-omics profile). In our OMC workflow and for each case, we first rank features based on their correlation with drug response and then train the k^{th} model with the top k features. The idea is to discard the many thousands of uninformative features by selecting the model trained only with those with the highest information content. The optimal k (k_{opt}), and thus a fast approximation to the optimal subset of features of size k_{opt} , corresponds to the model with the highest cross-validated performance in a given case. Results for all cases are reported on test sets that were neither used for training nor for model selection.

Model selection by stratified 5-fold cross-validation on the training set. 5-fold cross-validation was used as strategy to select the OMC. For our 5-fold cross-validation on the training set, we iteratively generated an internal training set (80% of n cell lines in training set) used to build the model and a validation set (remaining 20% of n cell lines in training set) to internally test. This ensured that our internal validation sets were of similar size to the external test set. To further guarantee the representativeness of the external test set, we used the same stratification procedure described in “Generation of non-overlapping training and test sets”. Per fold and per possible complexity between 2 and $n/2$, the correlation between predicted and observed log IC50 values was calculated using the Spearman correlation.

Complexities that failed to provide to outperform the minimal R_s threshold of 0.25 in any of the five validation folds were discarded. We also rejected models that predicted a constant output value regardless of input data. Enforcement of these rules slightly penalized lower complexities in scenarios where the response range (distance between minimal and maximal

log IC50) was low and the models predicted the same drug response irrespective of the biological features used input. The Optimal Model Complexity (OMC) was chosen as the complexity for which the median Spearman correlation across the five folds was highest. As such, it is the number of input features that results in the best median model performance across the five validation folds.

Pre-ranking feature input. The OMC strategy starts by first ranking the input features based on the strength of their association with the log IC50 for a particular case. For discrete inputs (SNV, CNA), this strength is calculated as the p-value of R's[59] `wilcox.test` function, as the corresponding Python implementation uses approximations. R objects were imported in Python using `rpy2` (<https://rpy2.bitbucket.io>). P-values for continuous input were calculated using a Spearman rank-correlation test[59].

Obtaining the OMC for a given drug fold. After ranking the features from likely to be most relevant (lowest p-value) to increasingly less significant, we defined that the optimal complexity be between 2 features and $n/2$ features, where n equals the number of cell lines available in the training fold for a particular case. The upper threshold ($n/2$ features) was chosen to guarantee that we have at least two data instances per feature. In combination with the feature reduction by OMC, the resulting models should be less prone to overfitting. We also trained models using the maximal possible complexity (i.e. all features), as some cases may require the combination of more features for optimal performance.

XGBoost (XGB). We downloaded the code for XGB (<https://github.com/dmlc/xgboost>). We first trained and tested all models using default parameters and all features to investigate if XGB captured some of the underlying biology and performed better than random. For the actual model complexity estimation, we opted to use a more conservative learning rate (0.05) than the default value of the implementation, a maximal tree depth of 6 and a larger number of trees (700), as described by Sheridan and co-workers[21]. Subsampling of row (cell line)

and column (feature) features per boosting iteration without replacement was set as 0.8 for both in the OMC estimation model. The 80% denotes the fraction of features and observations that was used to build a tree. Case-specific comprehensive hyperparameter tuning is expected to improve results further, but it is impractical as the number of cases is too large.

Final model performance estimation. For the construction of the final models, the full training set was used for model training. First, features were pre-ranked as described earlier. Second, the top OMC features were used as input for the models, which were again trained based on 10 random seeds. All other parameters were kept the same as in the random seed 5-fold CV loop (settings as in the conservative approach presented by Sheridan et al[21]). Finally, model performance was calculated using the Spearman metric between the ten predicted log IC50s and the corresponding test set values observed in the 10 cell lines of the external test set.

Pathway enrichment analysis. Enrichment analyses were performed using the Webgestalt tool[60] for the exploratory analyses for the case studies. A minimal significance of 0.05 with Benjamini-Hochberg correction over 1000 permutations for p-value estimation was set. The minimal threshold of features required in a category, for a category to be recognized as potentially interesting was defined as 5 (default tool setting). Only annotations for *Homo sapiens* were used for the GOSlim[61] and KEGG[62] pathway sets.

Statistical tests: Statistical tests were executed in R[59], using Python in combination with the rpy2 interface (<http://rpy.sourceforge.net/rpy2>).

DECLARATIONS

List of abbreviations

CCLE: Cancer Cell Line Encyclopedia

CTRP: Cancer Therapeutics Response Portal

GDSC: Genomics of Drug Sensitivity in Cancer

GO BP: Gene Ontology – Biological Process

OMC: Optimal Model Complexity

PDX: Patient-Derived Xenografts

RF: Random Forest

XGB: Extreme Gradient Boosting

Consent for publication

Not applicable.

Competing interests

The authors declare that they have no competing interests.

Availability of data and materials

Code supporting the conclusions of this paper and its instructions of use are available at

http://ballester.marseille.inserm.fr/GDSC_XGB-OMC.zip

The instructions also specify which input data files are used and hence have to be

downloaded from <https://www.cancerrxgene.org/>

REFERENCES

1. Neve RM, Chin K, Fridlyand J, Yeh J, Baehner FL, Fevr T, et al. A collection of breast cancer cell lines for the study of functionally distinct cancer subtypes. *Cancer Cell*. 2006;
2. Barretina J, Caponigro G, Stransky N, Venkatesan K, Margolin AA, Kim S, et al. The Cancer Cell Line Encyclopedia enables predictive modelling of anticancer drug sensitivity. *Nature* [Internet]. 2012 [cited 2012 Mar 29];483:603–307. Available from: <http://www.nature.com/nature/journal/v483/n7391/full/nature11003.html>
3. Rees MG, Seashore-Ludlow B, Cheah JH, Adams DJ, Price E V., Gill S, et al. Correlating chemical sensitivity and basal gene expression reveals mechanism of action. *Nat Chem Biol*.

2016;12:109–16.

4. Yang W, Soares J, Greninger P, Edelman EJ, Lightfoot H, Forbes S, et al. Genomics of Drug Sensitivity in Cancer (GDSC): a resource for therapeutic biomarker discovery in cancer cells. *Nucleic Acids Res* [Internet]. 2012 [cited 2012 Dec 10]; Available from: <http://nar.oxfordjournals.org/content/early/2012/11/22/nar.gks1111>

5. Garnett MJ, Edelman EJ, Heidorn SJ, Greenman CD, Dastur A, Lau KW, et al. Systematic identification of genomic markers of drug sensitivity in cancer cells. *Nature* [Internet]. 2012 [cited 2012 Mar 29];483:570–5. Available from: http://www.nature.com/nature/journal/v483/n7391/full/nature11005.html?WT.ec_id=NATURE-20120329

6. Iorio F, Knijnenburg TA, Vis DJ, Bignell GR, Menden MP, Schubert M, et al. A Landscape of Pharmacogenomic Interactions in Cancer. *Cell* [Internet]. Elsevier; 2016 [cited 2016 Aug 1];166:740–54. Available from: <http://linkinghub.elsevier.com/retrieve/pii/S0092867416307462>

7. Shoemaker RH. The NCI60 human tumour cell line anticancer drug screen. *Nat Rev Cancer* [Internet]. 2006;6:813–23. Available from: <http://www.nature.com/nrc/journal/v6/n10/full/nrc1951.html>

8. Ali M, Aittokallio T. Machine learning and feature selection for drug response prediction in precision oncology applications. *Biophys Rev* [Internet]. Springer Berlin Heidelberg; 2018 [cited 2018 Aug 21];1–9. Available from: <http://link.springer.com/10.1007/s12551-018-0446-z>

9. Azuaje F. Computational models for predicting drug responses in cancer research. *Brief Bioinform* [Internet]. Oxford University Press; 2016 [cited 2016 Oct 15];bbw065. Available from: <http://www.ncbi.nlm.nih.gov/pubmed/27444372>

10. De Niz C, Rahman R, Zhao X, Pal R. Algorithms for Drug Sensitivity Prediction.

Algorithms [Internet]. Multidisciplinary Digital Publishing Institute; 2016 [cited 2018 Jul 3];9:77. Available from: <http://www.mdpi.com/1999-4893/9/4/77>

11. Chen B, Butte AJ. Leveraging big data to transform target selection and drug discovery. *Clin Pharmacol Ther.* 2016;99:285–97.

12. Costello JC, Heiser LM, Georgii E, Gönen M, Menden MP, Wang NJ, et al. A community effort to assess and improve drug sensitivity prediction algorithms. *Nat Biotechnol* [Internet]. Nature Publishing Group, a division of Macmillan Publishers Limited. All Rights Reserved.; 2014 [cited 2014 Jul 14];32:1202–12. Available from: <http://dx.doi.org/10.1038/nbt.2877>

13. Ammad-ud-din M, Georgii E, Gönen M, Laitinen T, Kallioniemi O, Wennerberg K, et al. Integrative and personalized QSAR analysis in cancer by kernelized Bayesian matrix factorization. *J Chem Inf Model* [Internet]. American Chemical Society; 2014 [cited 2015 Sep 17];54:2347–59. Available from: <http://pubs.acs.org.gate2.inist.fr/doi/abs/10.1021/ci500152b>

14. Futreal PA, Coin L, Marshall M, Down T, Hubbard T, Wooster R, et al. A census of human cancer genes. *Nat Rev Cancer* [Internet]. Nature Publishing Group; 2004 [cited 2018 Aug 8];4:177–83. Available from: <http://www.nature.com/articles/nrc1299>

15. Naulaerts S, Dang CC, Ballester PJ, Naulaerts S, Dang CC, Ballester PJ, et al. Precision and recall oncology: combining multiple gene mutations for improved identification of drug-sensitive tumours. *Oncotarget* [Internet]. Impact Journals; 2017 [cited 2017 Oct 14];5. Available from: <http://www.oncotarget.com/fulltext/20923>

16. Menden MP, Iorio F, Garnett M, McDermott U, Benes CH, Ballester PJ, et al. Machine Learning Prediction of Cancer Cell Sensitivity to Drugs Based on Genomic and Chemical Properties. *PLoS One* [Internet]. 2013;8:e61318. Available from: <http://dx.doi.org/10.1371/journal.pone.0061318>

17. Riddick G, Song H, Ahn S, Walling J, Borges-Rivera D, Zhang W, et al. Predicting in vitro drug sensitivity using Random Forests. *Bioinformatics* [Internet]. 2011 [cited 2013 Oct 13];27:220–4. Available from: <http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3018816&tool=pmcentrez&rendertype=abstract>
18. Nguyen L, Dang CC, Ballester P, Nguyen L, Dang CC, Ballester P. Systematic assessment of multi-gene predictors of pan-cancer cell line sensitivity to drugs exploiting gene expression data. *F1000Research* [Internet]. 2016 [cited 2016 Dec 30];5:2927. Available from: <https://f1000research.com/articles/5-2927/v1>
19. Huang M, Shen A, Ding J, Geng M. Molecularly targeted cancer therapy: some lessons from the past decade. *Trends Pharmacol Sci* [Internet]. 2014 [cited 2015 Jul 7];35:41–50. Available from: <http://www.sciencedirect.com/science/article/pii/S0165614713002204>
20. Chen T, Guestrin C. XGBoost: A Scalable Tree Boosting System. *Proc 22nd ACM SIGKDD Int Conf Knowl Discov Data Min - KDD '16* [Internet]. New York, New York, USA: ACM Press; 2016 [cited 2018 Mar 21]. p. 785–94. Available from: <http://dl.acm.org/citation.cfm?doid=2939672.2939785>
21. Sheridan RP, Wang WM, Liaw A, Ma J, Gifford EM. Extreme Gradient Boosting as a Method for Quantitative Structure-Activity Relationships. *J Chem Inf Model*. 2016;56:2353–60.
22. He T, Heidemeyer M, Ban F, Cherkasov A, Ester M. SimBoost: a read-across approach for predicting drug-target binding affinities using gradient boosting machines. *J Cheminform*. 2017;9.
23. Haury A-C, Gestraud P, Vert J-P. The Influence of Feature Selection Methods on Accuracy, Stability and Interpretability of Molecular Signatures. Teh M-T, editor. *PLoS One* [Internet]. Public Library of Science; 2011 [cited 2017 Jun 6];6:e28210. Available from:

<http://dx.plos.org/10.1371/journal.pone.0028210>

24. Lee S-I, Celik S, Logsdon BA, Lundberg SM, Martins TJ, Oehler VG, et al. A machine learning approach to integrate big data for precision medicine in acute myeloid leukemia. *Nat Commun.* 2018;9:42.
25. Dang CC, Peón A, Ballester PJ. Unearthing new genomic markers of drug response by improved measurement of discriminative power. *BMC Med Genomics.* 2018;11.
26. Alzubaidi A, Cosma G. A multivariate feature selection framework for high dimensional biomedical data classification. 2017 IEEE Conf Comput Intell Bioinforma Comput Biol [Internet]. IEEE; 2017 [cited 2018 Sep 3]. p. 1–8. Available from: <http://ieeexplore.ieee.org/document/8058528/>
27. Schwarze K, Buchanan J, Taylor JC, Wordsworth S. Are whole-exome and whole-genome sequencing approaches cost-effective? A systematic review of the literature. *Genet Med.* 2018;
28. Poirier JT, Gardner EE, Connis N, Moreira AL, De Stanchina E, Hann CL, et al. DNA methylation in small cell lung cancer defines distinct disease subtypes and correlates with high expression of EZH2. *Oncogene.* 2015;34:5869–78.
29. Codignola A, Tarroni P, Clementi F, Pollo A, Lovallo M, Carbone E, et al. Calcium channel subtypes controlling serotonin release from human small cell lung carcinoma cell lines. *J Biol Chem.* 1993;268:26240–7.
30. Taniwaki M, Daigo Y, Ishikawa N, Takano A, Tsunoda T, Yasui W, et al. Gene expression profiles of small-cell lung cancers: Molecular signatures of lung cancer. *Int J Oncol.* 2006;29:567–75.
31. Hubaux R, Thu KL, Coe BP, Macaulay C, Lam S, Lam WL. EZH2 promotes E2F-driven SCLC tumorigenesis through modulation of apoptosis and cell-cycle regulation. *J Thorac Oncol.* 2013;

32. Bilir B, Kucuk O, Moreno CS. Wnt signaling blockage inhibits cell proliferation and migration, and induces apoptosis in triple-negative breast cancer cells. *J Transl Med*. 2013;11.
33. Shen C-Y, Chen L-H, Lin Y-F, Lai L-C, Chuang EY, Tsai M-H, et al. Mitomycin C treatment induces resistance and enhanced migration via phosphorylated Akt in aggressive lung cancer cells. *Oncotarget*. 2016;5.
34. He B, Santamaria R, Xu W, Cols M, Chen K, Puga I, et al. The transmembrane activator TACI triggers immunoglobulin class switching by activating B cells through the adaptor MyD88. *Nat Immunol*. 2010;
35. Kang YH, Lee KA, Ryu CJ, Lee HG, Lim JS, Park SN, et al. Mitomycin C induces apoptosis via Fas/FasL dependent pathway and suppression of IL-18 in cervical carcinoma cells. *Cancer Lett*. 2006;237:33–44.
36. Siegel D, Yan C, Ross D. NAD(P)H:quinone oxidoreductase 1 (NQO1) in the sensitivity and resistance to antitumor quinones. *Biochem. Pharmacol*. 2012. p. 1033–40.
37. Shkreta L, Chabot B. The RNA splicing response to DNA damage. *Biomolecules*. 2015. p. 2935–77.
38. Al-Saffar NMS, Jackson LE, Raynaud FI, Clarke PA, De Molina AR, Lacal JC, et al. The phosphoinositide 3-kinase inhibitor PI-103 downregulates choline kinase α leading to phosphocholine and total choline decrease detected by magnetic resonance spectroscopy. *Cancer Res*. 2010;70:5507–17.
39. Prevo R, Deutsch E, Sampson O, Diplexcito J, Cenge K, Harper J, et al. Class I PI3 kinase inhibition by the pyridinylfuranopyrimidine inhibitor PI-103 enhances tumor radiosensitivity. *Cancer Res*. 2008;68:5915–23.
40. Raynaud FI, Eccles S, Clarke PA, Hayes A, Nutley B, Alix S, et al. Pharmacologic characterization of a potent inhibitor of class I phosphatidylinositide 3-kinases. *Cancer Res*.

2007;67:5840–50.

41. Szklarczyk D, Franceschini A, Wyder S, Forslund K, Heller D, Huerta-Cepas J, et al. STRING v10: Protein-protein interaction networks, integrated over the tree of life. *Nucleic Acids Res.* 2015;
42. Kojima K, Shimanuki M, Shikami M, Samudio IJ, Ruvolo V, Corn P, et al. The dual PI3 kinase/mTOR inhibitor PI-103 prevents p53 induction by Mdm2 inhibition but enhances p53-mediated mitochondrial apoptosis in p53 wild-type AML. *Leukemia.* 2008;22:1728–36.
43. Keisner S V., Shah SR. Pazopanib: The newest tyrosine kinase inhibitor for the treatment of advanced or metastatic renal cell carcinoma. *Drugs.* 2011. p. 443–54.
44. Fruehauf JP, Alger B, Parmakhtiar B, Jakowatz JG, Bettis C, Chuang T, et al. A phase II single arm study of pazopanib and paclitaxel as first-line treatment for unresectable stage III and stage IV melanoma: Interim analysis. *ASCO Meet Abstr .* 2012;30:8524.
45. Kumar R, Knick VB, Rudolph SK, Johnson JH, Crosby RM, Crouthamel M-C, et al. Pharmacokinetic-pharmacodynamic correlation from mouse to human with pazopanib, a multikinase angiogenesis inhibitor with potent antitumor and antiangiogenic activity. *Mol Cancer Ther.* 2007;6:2012–21.
46. Hurvitz SA, Shatsky R, Harbeck N. Afatinib in the treatment of breast cancer. *Expert Opin Investig Drugs.* 2014;
47. Connell CM, Doherty GJ. Activating HER2 mutations as emerging targets in multiple solid cancers. *ESMO Open.* 2017.
48. Nakai K, Hung MC, Yamaguchi H. A perspective on anti-EGFR therapies targeting triple-negative breast cancer. *Am. J. Cancer Res.* 2016.
49. Vakiani E, Solit DB. KRAS and BRAF: Drug targets and predictive biomarkers. *J. Pathol.* 2011.
50. Ciombor KK, Bekaii-Saab T. Selumetinib for the treatment of cancer. *Expert Opin*

Investig Drugs. Taylor & Francis; 2015;24:111–23.

51. Gross AM, Wolters P, Baldwin A, Dombi E, Fisher MJ, Weiss BD, et al. SPRINT: Phase II study of the MEK 1/2 inhibitor selumetinib (AZD6244, ARRY-142886) in children with neurofibromatosis type 1 (NF1) and inoperable plexiform neurofibromas (PN). *J Clin Oncol*. American Society of Clinical Oncology; 2018;36:10503.

52. Goodspeed A, Heiser LM, Gray JW, Costello JC. Tumor-derived Cell Lines as Molecular Models of Cancer Pharmacogenomics. *Mol Cancer Res* [Internet]. 2015 [cited 2015 Nov 8];1541-7786.MCR-15-0189-. Available from:

<http://mcr.aacrjournals.org/content/early/2015/08/06/1541-7786.MCR-15-0189.abstract>

53. Gao H, Korn JM, Ferretti S, Monahan JE, Wang Y, Singh M, et al. High-throughput screening using patient-derived tumor xenografts to predict clinical trial drug response. *Nat Med* [Internet]. Nature Publishing Group; 2015 [cited 2015 Oct 21];21:1318–25. Available from: <http://www.nature.com/gate2.inist.fr/nm/journal/v21/n11/full/nm.3954.html>

54. Domcke S, Sinha R, Levine DA, Sander C, Schultz N. Evaluating cell lines as tumour models by comparison of genomic profiles. *Nat Commun*. 2013;4.


55. Vincent KM, Findlay SD, Postovit LM. Assessing breast cancer cell lines as tumour models by comparison of mRNA expression profiles. *Breast Cancer Res*. 2015;17.

56. Geeleher P, Zhang Z, Wang F, Gruener RF, Nath A, Morrison G, et al. Discovering novel pharmacogenomic biomarkers by imputing drug response in cancer patients from large genomics studies. *Genome Res*. 2017;27:1743–51.

57. Preussner J, Bayer J, Kuenne C, Looso M. ADMIRE: Analysis and visualization of differential methylation in genomic regions using the Infinium HumanMethylation450 Assay. *Epigenetics and Chromatin*. 2015;8.

58. Smedley D, Haider S, Ballester B, Holland R, London D, Thorisson G, et al. BioMart--biological queries made easy. *BMC Genomics*. 2009;10:22.

- 1 59. R Development Core Team R. R: A Language and Environment for Statistical
2 Computing. R Found. Stat. Comput. 2011.
3
4 60. Wang J, Duncan D, Shi Z, Zhang B. WEB-based GEne SeT AnaLysis Toolkit
5 (WebGestalt): update 2013. Nucleic Acids Res. 2013;41.
6
7 61. The Gene Ontology Consortium. Expansion of the Gene Ontology knowledgebase and
8 resources. Nucleic Acids Res. 2017;
9
10 62. Kanehisa M, Furumichi M, Tanabe M, Sato Y, Morishima K. KEGG: New perspectives
11 on genomes, pathways, diseases and drugs. Nucleic Acids Res. 2017;
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65



[Click here to access/download](#)

Supplementary Material

Naulaerts et al XGB-OMC GDSC-SI_final.docx



Date: 24/January/2019

Dear Editor,

We are pleased to submit our manuscript for consideration at GigaScience. In this letter, I would like to tell you about the significance of our work for Personalised Oncology.

Personalised oncology has largely been based on analysing pharmacogenomics data to identify a specific somatic mutation able to discriminate between responders and non-responders. Multiple initiatives, such as the French Plan for Genomic Medicine 2025 [1], herald the arrival of a much larger volume of cancer pharmacogenomics data. State-of-the-art technological platforms will carry out the standard analysis of these data sets, thus discovering new genomic markers for some more cases (i.e. drug-cancer type pairs). However, this strongly reductionist approach will continue to suffer from the same limitations [2]: many cases will remain without a marker as a single gene will not be sufficiently predictive of tumour response and the many responders without the actionable mutation will still be missed by the marker where one is found. Consequently, personalised oncology will continue to benefit only a very small proportion of patients [2]. The need for better methodologies is reflected in this crucial statement [1] with which we particularly agree: *"In addition to servers, cores and intensive calculation methods, the development of software for data mining and modelling will make it possible to meet the challenge of genomic medicine and, more broadly, personalized medicine."*

There are several reasons why we are employing cancer cell lines with this purpose. Cell lines have demonstrated to be highly useful cancer models, despite their well-known limitations. Given their suitability for high-throughput experiments, they are also the model for which most phenotypic and pharmacological information is publicly available. For many drug-cancer type pairs, only much smaller amounts of pharmacogenomics data sets, if any, has been released using more patient-relevant *in vivo* models. Consequently, often the only data that can be used to predict drug sensitivity in these cases comes from cell lines, thus having their niche in guiding precision oncology efforts.

Our study integrates state-of-the-art nonlinear machine learning algorithms (XGBoost; XGB for short) with a new feature selection strategy (OMC) to generate very predictive models using only concise sets of cell line features. The latter means that only a fraction of the initially considered features need to be determined in additional cell lines, which allows better interpretability and less resource-intensive application to other tumour samples. Here we identify the most predictive molecular profile for each of drug-cancer type. XGB-OMC predictors exploiting gene expression or DNA methylation profiles are particularly innovative, as these profiles have not been previously analysed with nonlinear regression

algorithms able to find a concise subset of predictive features. Overall, we have found 118 predictive models for 62 drugs across five cancer types using four molecular profiles, which provide new insights into this key problem.

Thanks again for your time to consider our submission.

Yours sincerely,



Dr Pedro BALLESTER, on behalf of the co-authors
Group leader at the CRCM
INSERM research scientist

- [1] F Lethimonnier, Y Levy (2018) "Genomic Medicine France 2025". *Annals of Oncology*, mdy027.
[2] Huang, M., Shen, A., Ding, J. & Geng, M. (2014) "Molecularly targeted cancer therapy: some lessons from the past decade". *Trends Pharmacol. Sci.* **35**, 41–50.